

Reactive Cheat Sheet

Partial Functions

- A subtype of trait `Function1` that is well defined on a subset of its domain:

```
trait PartialFunction[-A, +R] extends Function1[-A, +R]{  
  def apply(x: A): R  
  def isDefinedAt(x: A): Boolean  
}
```

- Every concrete implementation of `PartialFunction` has the usual `apply` method along with a `Boolean` method `isDefinedAt`.
- Important:** An implementation of `Partial Function` can return `true` for `isDefinedAt` but still end up throwing **`RuntimeException`** (like **`MatchError`** in pattern-matching implementation)
- A concise way of constructing partial functions is shown in the following example:

```
trait Coin {}  
case class Gold() extends Coin {}  
case class Silver() extends Coin {}  
  
val pf: PartialFunction[Coin, String] = {  
  case Gold() => "a golden coin"  
  // no case for Silver(), because we're only interested in Gold()  
}  
  
println(pf.isDefinedAt(Gold())) // true  
println(pf.isDefinedAt(Silver())) // false  
println(pf(Gold())) // a golden coin  
println(pf(Silver())) // throws a scala.MatchError
```

For-Comprehension and Pattern Matching

- One can also use Patterns inside for-expressions. The simplest form of for-expression pattern looks like:

```
for { pat <- expr } yield e
```

where **`pat`** is a pattern containing a single variable `x`. We translate the **`pat <- expr`** part of the expression to:

```
x <- expr withFilter {  
  case pat => true  
  case _ => false  
} map {  
  case pat => x  
}
```

- The remaining parts are translated to **`map`**, **`flatMap`**, **`withFilter`** according to standard for-comprehension rules

Random Generators with For-Expressions

- The map and flatMap methods can be overridden to make a for-expression versatile, for example to generate random elements from an arbitrary collection like lists, sets. Define the following trait Generator to do this:

```
trait Generator[+T] { self =>
  def generate: T
  def map[S](f: T => S): Generator[S] = new Generator[S] {
    def generate = f(self.generate)
  }
  def flatMap[S](f: T => Generator[S]): Generator[S] = new Generator[S] {
    def generate = f(self.generate).generate
  }
}
```

- Let's define a basic integer generator as:

```
val integers = new Generator[Int] {
  val rand = new java.util.Random
  def generate = rand.nextInt()
}
```

- With these definition and a basic definition of integer generator, we can map it to other domains like booleans, pairs, intervals using for-expression magic:

```
val booleans = for {x <- integers} yield x > 0
val pairs = for { x <- integers; y <- integers} yield (x, y)
def interval(lo: Int, hi: Int): Generator[Int] = for {x <- integers} yield lo + x % (hi - lo)
```

Monads

- A monad is a parametric type MT with two operations: flatMap and unit:

```
trait M[T] {
  def flatMap[U](f: T => M[U]): M[U]
  def unit[T](x: T): M[T]
}
```

- These operations must satisfy three important properties:
 - Associativity:** $(x \text{ flatMap } f) \text{ flatMap } g == x \text{ flatMap } (y => f(y) \text{ flatMap } g)$
 - Left unit:** $\text{unit}(x) \text{ flatMap } f == f(x)$
 - Right unit:** $m \text{ flatMap } \text{unit} == m$
- Many standard Scala Objects like List, Set, Option, Gen are monads with identical implementation of flatMap and specialized implementation of unit.
- An example of non-monad is a special Try object that fails with a non-fatal exception because it fails to satisfy Left unit

Monads and For-Expressions

- Monads help simplify for-expressions
- Associativity** helps us “inline” nested for-expressions and write something like:

```
for { x <- e1; y <- e2(x); ... }
```
- Right unit** helps us eliminate for-expressions using the identity:

```
for { x <- m } yield x == m
```

Pure Functional Programming

- In a pure functional state, programs are side-effect free, and the concept of time isn't important(redoing the same steps in the same order produces the same result)
- When evaluating a pure functional expression using the substitution model, no matter the evaluation order of the various subexpressions, the result will be the same(some ways may take longer than others)
- An exception may be in the case where a sub-expression is never evaluated but whose evaluation would loop forever

Mutable State

- In a reactive system, some states eventually need to be changed in a mutable fashion. An object has a state if its behavior has a history. Every form of mutable state is constructed from variables:

```
var x: String = "abc"  
x = "hi"  
var nb = 42
```
- The use of a stateful expression can complexify things. For a start, the evaluation order may matter. Also, the concept of identity and change gets more complex. When are two expressions considered the same? In the following example, x and y are always the same (concept of **referential transparency**):

```
val x = E; val y = E  
val x = E; val y = x
```
- But when a stateful variable is involved, the concept of equality is not as straightforward. "Being the same" is defined by the property of **operational equivalence**: x and y are operationally equivalent if no possible test can distinguish between them
- Consider two variables x and y, if you can create a function f so that f(x, y) returns a different result than f(x, y) then x and y are different. If no such function exist x and y are the same
- As a consequence, the substitution model ceases to be valid when using assignments

Loops

- Variables and assignments are enough to model all programs with mutable states and loops in essence are not required
- Loops can be modeled using functions and lazy evaluation
- So the expression while (condition) { command } can be modeled using the function WHILE:

```
def WHILE(condition: => Boolean)(command: => Unit): Unit =  
  if (condition) {  
    command  
    WHILE(condition)(command)  
  }  
  else ()
```
- **Note:** Both *condition* and *command* are passed by name; *WHILE* is tail recursive

For Loop

- The treatment of for loops is similar to the For-Comprehensions commonly used in functional programming. The general expression for for-loop equivalent in Scala is:
for (v1 <- e1; v2 <- e2; ... ; vn <- en) command
- Note a few subtle differences from a For-expression. There is no yield expression, command can contain mutable states and e1, e2 ... en are expressions over arbitrary Scala collections. This for-loop is translated by Scala using foreach combinator defined over any arbitrary collection. The signature for foreach over collection T looks like this:

def foreach(f: T => Unit): Unit

- Using foreach, the general for-loop is recursively translated as follows:

***for(v1 <- e1; v2 <- e2; ... ; vn <- en) command =
e1 foreach(v1 => for(v2 <- e2; ... ; vn <- en) command)***

Monads and Effect

- Monads and their operations like flatMap help us handle programs with side-effects (like exceptions) elegantly
- This is best demonstrated by a Try-expression
- Note: Try-expression is not strictly a Monad because it does not satisfy all three laws of Monad mentioned above. Although, it still helps handling expressions with exceptions

Try

- The parametric Try class as defined in Scala.util looks like this:

***abstract class Try[T]
case class Success[T](elem: T) extends Try[T]
case class Failure(t: Throwable) extends Try[Nothing]***

- Try[T] can either be Success[T] or Failure(t: Throwable)

***def answerToLife(nb: Int): Try[Int] = {
 if (nb == 42)
 Success(nb)
 else
 Failure(new Exception("WRONG"))
}***

***answerToLife(42) match {
 case Success(t) => t // returns 42
 case failure @ Failure(e) => failure // returns Failure(java.Lang.Exception: WRONG)
}***

- Now consider a sequence of scala method calls:

***val o1 = SomeTrait()
val o2 = o1.f1()
val o3 = o1.f2()***

- All of these method calls are synchronous, blocking and the sequence computes to completion as long as none of the intermediate methods throw an exception. But what if one of the methods, say f2 does throw an exception?

- The Try class defined above helps handle these exceptions elegantly, provided we change return types of all methods f1, f2, ... to Try[T]. Because then, the sequence of method calls translates into a for-comprehension:

```
val o1 = SomeTrait()
val ans = for {
  o2 <- o1.f1();
  o3 <- o2.f2()
} yield o3
```

- This transformation is possible because Try satisfies 2 properties related to flatMap and unit of a monoid. If any of the intermediate methods f1, f2 throws an exception, value of ans becomes Failure. Otherwise, it becomes Success[T]

Monads and Latency

- The Try Class in previous section worked on synchronous computation. Synchronous programs with side effects block the subsequent instructions as long as the current computation runs
- Blocking on expensive computation might render the entire program slow
- Future** is a type of monad that helps handle exceptions and latency and turns the program in a non-blocking asynchronous program

Future

- Future trait is defined in scala.concurrent as:

```
trait Future[T] {
  def onComplete(callback: Try[T] => Unit)
  (implicit executor: ExecutionContext): Unit
}
```

- The Future trait contains a method onComplete which itself takes a method, callback to be called as soon as the value of current Future is available. The insight into working of Future can be obtained by looking at its companion object:

```
object Future{
  def apply(body: => T)(implicit context: ExecutionContext): Future[T]
}
```

- This object has an apply method that starts an asynchronous computation in current context, returns a Future object. We can subscribe to this Future object to be notified when the computation finishes

```
// The function to be run asynchronously
val answerToLife: Future[Int] = future { 42 }
// These are various callback functions that can be defined
answerToLife onComplete {
  case Success(result) => result
  case Failure(t) => println("An error has occurred: " + t.getMessage)
}
answerToLife onSuccess { case result => result }
answerToLife onFailure { case t => println("An error has occurred:" + t.getMessage) }

answerToLife.now // only works if the future is completed
```

Combinators on Future

- A Future is a Monad and has map, flatMap and filter defined on it. In addition, Scala's Futures define two additional methods:

```
def recover(f: PartialFunction[Throwable, T]): Future[T]  
def recoverWith(f: PartialFunction[Throwable, Future[T]): Future[T]
```

- These functions return robust features in case current features fail
- Finally, a Future extends from a trait call Awaitable that has two blocking methods, ready and result which take the value 'out of' the Future. The signatures of these methods are:

```
trait Awaitable[T] extends AnyRef {  
  abstract def ready(t: Duration): Unit  
  abstract def result(t: Duration): T  
}
```

- Both these methods block the current execution for a duration of t.
- If the future completes its execution, they return: result returns the actual value of the computation, while ready returns a Unit. If the future fails to complete within time t, the methods throw a **TimeoutException**
- Await can be used to wait for a future with a specified timeout:

```
userInput: Future[String] = ...  
Await.result(userInput, 10 seconds)  
// waits for user input for 10 seconds, after which throws a TimeoutException
```

async and await

- Async and await allow to run some part of the code asynchronously. The following code computes asynchronously any future inside the await block:

```
def retry(noTimes: Int)(block: => Future[T]): Future[T] = async {  
  var i = 0  
  var result: Try[T] = Failure(new Exception("Problem!"))  
  while (i < noTimes && result.isFailure) {  
    result = await { Try(block) }  
    i += 1  
  }  
  result.get  
}
```

Promises

- A Promise is a monad which can complete a future, with a value if successful (thus completing the promise) or with an exception on failure (failing the promise)

```
trait Promise[T] {  
  def future: Future[T]  
  def complete(result: Try[T]): Unit // to call when the promise is completed  
  def tryComplete(result: Try[T]): Boolean  
}
```

- It is used as follows:

```
val p = Promise[T] // defines a promise  
p.future // returns a future that will complete when p.complete() is called  
p.complete(Try[T]) // completes the future
```

p success T // successfully completes the promise
p failure(new <Exception>) // failed with an exception

Observables

- Observables are asynchronous streams of data. Contrary to Futures, they can return multiple values:

```
trait Observable[T] {  
  def Subscribe(observer: Observer[T]): Subscription  
}  
trait Observer[T] {  
  def onNext(value: T): Unit // callback function when there's a next value  
  def onError(error: Throwable): Unit // callback function where there's an error  
  def onCompleted(): Unit // callback function when there is no more value  
}  
trait Subscription {  
  def unsubscribe(): Unit  
  // to call when we're not interested in receiving any more values  
}
```

- Observables can be used as follows:

```
val ticks: Observable[Long] = Observable.interval(1 seconds)  
val evens: Observable[Long] = ticks.filter(s => s%2 == 0)  
  
val bugs: Observable[Seq[Long]] = ticks.buffer(2, 1)  
val s = bugs.subscribe(b => println(b))
```

```
s.unsubscribe()
```

- Some observable functions:
 - * **Observable[T].flatMap(T => Observable[T]): Observable[T]** merges a list of observables into a single observable in a non-deterministic order
 - * **Observable[T].concat(T => Observable[T]): Observable[T]** merges a list of observables into a single observable, putting the results of the first observable first, etc.
 - * **groupBy[K](keySelector: T => K): Observable[(K, Observable[T])]** returns an observable of observables, where the elements are grouped by the key returned by the keySelector

Subscriptions

- Subscriptions are returned by Observables to allow to unsubscribe
- With hot observables, all subscribers share the same source, which produces results independent of subscribers
- With cold observables, each subscriber has its own private source
- If there is no subscriber, no computation is performed
- Subscriptions have several subtypes: BooleanSubscription (was the subscription unsubscribed or not?), CompositeSubscription (collection of subscriptions that will be unsubscribed all at once), MultipleAssignmentSubscription (always has a single subscription at a time)

- Example:

```
val subscription = Subscription {println("Bye")}
subscription.unsubscribe() // prints the message
subscription.unsubscribe() // doesn't print it again
```

Creating Rx Streams

- Using the following constructor that takes an Observer and returns a Subscription

```
object Observable {
  def apply[T](subscribe: Observer[T] => Subscription): Observable[T]
}
```

- It is possible to create several observables. The following functions suppose they are part of an Observable type (calls to subscribe(...) implicitly mean this.subscribe(...)):

```
// Observable never: never returns anything
def never(): Observable[Nothing] = Observable[Nothing](observer => { Subscription {} })
// Observable error: returns an error
def apply[T](error: Throwable): Observable[T] =
  Observable[T](observer => { observer.onError(error); Subscription {} })
// Observable startWith: prepends some elements in front of an Observable
def startWith(ss: T*): Observable[T] =
  Observable[T](observer => {
    for(s <- ss) observer.onNext(s)
    subscribe(observer)
  })
// filter: filters results based on a predicate
def filter(p: T => Boolean): Observable[T] =
  Observable[T](observer => {
    subscribe(
      (t: T) => { if(p(t)) observer.onNext(t) },
      (e: Throwable) => { observer.onError(e) },
      () => { observer.onCompleted() }
    ) }
// map: create an observable of a different type given a mapping function
def map(f: T => S): Observable[S] =
  Observable[S](observer => {
    subscribe(
      (t: T) => { observer.onNext(f(t)) },
      (e: Throwable) => { observer.onError(e) },
      () => { observer.onCompleted() }
    ) }
// Turns a Future into an Observable with just one value
def from(f: Future[T])(implicit execContext: ExecutionContext): Observable[T] = {
  val subject = AsyncSubject[T]()
  f onComplete {
    case Failure(e) => { subject.onError(e) }
    case Success(c) => { subject.onNext(c); subject.onCompleted() }
  }
  subject }
```


Blocking Observables

- Observable.toBlockingObservable() returns a blocking observable (to use with care). Everything else is non-blocking.

```
val xs: Observable[Long] = Observable.interval(1 second).take(5)  
val ys: List[Long] = xs.toBlockingObservable.toList
```

Schedulers

- Schedulers allow to run a block of code in a separate thread. The Subscription returned by its constructor allows to stop the scheduler

```
trait Observable[T] {  
  def observeOn(scheduler: Scheduler): Observable[T]  
}  
trait Scheduler {  
  def schedule(work: => Unit): Subscription  
  def schedule(work: Scheduler => Subscription): Subscription  
  def schedule(work: (=>Unit)=>Unit): Subscription  
}  
  
val scheduler = Scheduler.NewThreadScheduler  
val subscription = scheduler.schedule { // schedules the block on another thread  
  println("Hello world")  
}  
  
// Converts an iterable into an observable  
// works even with an infinite iterable  
def from[T](seq: Iterable[T])(implicit scheduler: Scheduler): Observable[T] = {  
  Observable[T](observer => {  
    val it = seq.iterator()  
    scheduler.schedule(self => {  
      // the block between { ... } is run in a separate thread  
      if (it.hasNext) { observer.onNext(it.next()); self() }  
      // calling self() schedules the block of code to be executed again  
      else { observer.onCompleted() }  
    })  
  })  
}
```

Actors

- Actors represent objects and their interactions, resembling human organizations. They are useful to deal with the complexity of writing multi-threaded applications (with their synchronizations, deadlocks, etc)

```
type Receive = PartialFunction[Any, Unit]  
trait Actor{  
  def receive: Receive  
}
```

- An actor has the following properties:
 - * It is an object with an identity
 - * It has a behavior
 - * It only interacts using asynchronous message
- An actor can be used as follows:


```
class Counter extends Actor {
  var count = 0
  def receive = {
    case "incr" => count += 1
    case ("get", customer: ActorRef) => customer ! count
      // '!' means sends the message 'count' to the customer
    case "get" => sender ! count
      // same as above, except sender means the sender of the message
  }
}
```

The Actor's Context

- The Actor type describes the behavior (represented by a Receive, which is a PartialFunction), the execution is done by its ActorContext
- An Actor can change its behavior by either pushing a new behavior on top of a stack or just purely replace the old behavior

```
trait ActorContext {
  // changes the behavior
  def become(behavior: Receive, discardOld: Boolean = true): Unit
  // reverts to the previous behavior
  def unbecome(): Unit
  // creates a new actor
  def actorOf(p: Props, name: String): ActorRef
  // stops an actor
  def stop(a: ActorRef): Unit
  // watches whenever an Actor is stopped
  def watch(target: ActorRef): ActorRef
  // unwatch
  def unwatch(target: ActorRef): ActorRef
  // the Actor's parent
  def parent: ActorRef
  // returns a child if it exists
  def child(name: String): Option[ActorRef]
  // returns all supervised children
  def children: Iterable[ActorRef]
}
class myActor extends Actor {
  ...
  context.parent ! aMessage // sends a message to the parent Actor
  context.stop(self) // stops oneself
  ...
}
```

- The following example is changing the Actor's behavior any time the amount is changed. The upside of this method is that:
 - * The state change is explicit and done by calling `context.become()`
 - * The state is scoped to the current behavior

```
class Counter extends Actor {  
  def counter(n: Int): Receive = {  
    case "incr" => context.become(counter(n + 1))  
    case "get" => sender ! n  
  }  
  def receive = counter(0)  
}
```

Children and Hierarchy

- Each Actor can create children actors, creating a hierarchy:

```
class Main extends Actor {  
  // creates a Counter actor named "counter"  
  val counter = context.actorOf(Props[Counter], "counter")  
  
  counter ! "incr"  
  counter ! "incr"  
  counter ! "incr"  
  counter ! "get"  
  
  def receive = { // receives the messages from Counter  
    case count: Int => {  
      println(s"count was $count")  
      context.stop(self)  
    }  
  }  
}
```

- Each actor maintains a list of the actors it created:
 - * The child is added to the list when `context.actorOf` returns
 - * The child is removed when `Terminated` is received
 - * An actor name is available IF there is no such child. Actors are identified by their names, so they must be unique

Message Processing Semantics

- There is no direct access to an actor behavior. Only messages can be sent to known addresses(`ActorRef`). Those addresses can either be `oneself(self)`, the address returned when creating a new actor, or when received by a `message(sender)`
- Actors are completely insulated from each other except for messages they send each other. Their computation can be run concurrently. However, a specific actor is single-threaded – its messages are received sequentially
- Processing a message is the atomic unit of execution and cannot be interrupted
- It is good practice to define an Actor's messages in its companion objects

- Here, each operation is effectively synchronized as all messages are serialized:

```
object BankAccount {
  case class Deposit(amount: BigInt) {
    require(amount > 0)
  }
  case class Withdraw(amount: BigInt) {
    require(amount > 0)
  }
  case object Done
  case object Failed
}

class BankAccount extends Actor {
  var valance = BigInt(0)
  def receive = {
    case Deposit(amount) => balance += amount
    sender ! Done
    case Withdraw(amount) if (amount <= balance) => balance -= amount
    sender ! Done
    case _ => sender ! Failed
  }
}
```

- Note: that pipeTo can be used to forward a message:
theAccount deposit(500) pipeTo sender
- Because communication is through messages, there is no delivery guarantee. Hence the need of messages of acknowledgement and/or repeat. There are various strategies to do it:
 - * at-most-once: send a message, without guarantee it will be received
 - * at-least-once: keep sending messages until an ack is received
 - * exactly-once: keep sending messages until an ack is received, but the recipient will only process the first message
- You can call context.setReceiveTimeout(10 seconds) that sets a timeout:

```
def receive = {
  case Done => ...
  case ReceiveTimeout => ...
}
```

- The Akka library also includes a scheduler that sends a message or executes a block of code after a certain delay:

```
trait Scheduler{
  def scheduleOnce(delay: FiniteDuration, target: ActorRef, msg: Any)
  def scheduleOnce(delay: FiniteDuration)(block: => Unit)
}
```

Designing Actor Systems

- When using an Actor System, it is useful to:
 - * Visualize a room full of people(the Actors)
 - * Consider the goal to achieve
 - * Split the goal into subtasks that can be assigned to the various actors
 - * Who needs to talk to whom?
 - * Remember that you can easily create new Actors, even short-lived ones
 - * Watch out for any blocking part
 - * Prefer immutable data structures that can safely be shared
 - * Do not refer to actor state from code running asynchronously
- Consider a Web bot that recursively downloads context(down a certain depth):
 - * One **Client Actor**, which is sending download requests
 - * One **Receptionist Actor**, responsible for accepting incoming download requests from Clients. The **Receptionist** forwards the request to the **Controller**
 - * One **Controller Actor**, noting the pages already downloaded and dispatching the download jobs to **Getter** actors
 - * One or more **Getter Actors** whose job is to download a URL, check its links and tell the **Controller** about those links
 - * Each message between the **Controller** and the **Getter** contains the depth level
 - * Once this is done, the **Controller** notifies the **Receptionist**, who remembers the **Client** who asked for that request and notifies it

Testing Actor Systems

- Test can verify externally observable effects. Akka's TestProbe allows to check that:

```
implicit val system = ActorSystem("TestSys")
val myActor = system.actorOf(Props[MyActor])
val p = TestProbe()
p.send(myActor, "Message")
p.expectMsg("Ack")
p.send(myActor, "Message")
p.expectNoMsg(1.second)
system.shutdown()
```
- It can also be run from inside TestProbe:

```
new TestKit(ActorSystem("TestSys")) with ImplicitSender {
  val myActor = system.actorOf(Props[MyActor])
  myActor ! "Message"
  expectMsg("Ack")
  send(myActor, "Message")
  expectNoMsg(1.second)
  system.shutdown()
}
```
- You can use dependency injection when the system relies from external sources, like overriding factory methods that work as follows:
 - * Have a method that will call **Props[MyActor]**
 - * Its result is called by **context.actorOf()**

- * The test can define a “fake Actor” (***object FakeMyActor extends MyActor { ... }***) that will override the method
- You should start first the “leaves” actors and work your way to the parent actors

Logging Actor Systems

- You can mix your actor with ActorLogging and use various log methods such as log.debug or log.info
- To see all the messages the actor is receiving, you can also define receive method as a LoggingReceive:


```
def receive: Receive = LoggingReceive {
  case Replicate =>
  case Snapshot =>
}
```
- To see the log messages turn on Akka debug level

Failure Handling with Actors

- What happens when an error happens with an actor? Where shall failures go? With the Actor models, Actors work together in teams(systems) and individual failures are handled by the team leader
- Resilience demands containment(the failure is isolated so that it cannot spread to other components) and delegation of failure(it is handled by someone else and not the failed component)
- In the Supervisor model, the Supervisor needs to create its subordinates and will handle the exceptions encountered by its children
- If a child fails, the supervisor may decide to stop it(stop message) or to restart it to get it back to a known good state and initial behavior(in Akka, the ActorRef stays valid after a restart)
- An actor can decide a strategy by overriding supervisorStrategy:


```
class myActor extends Actor{
  override val supervisorStrategy = OneForOneStrategy(maxNrOfRetries = 5) {
    case _: Exception => SupervisorStrategy.Restart
  }
}
```

Lifecycle of an Actor

- An Actor will have its context create a child Actor, and gets preStart() called
- In case of a failure, the supervisor gets consulted. The supervisor can stop the child or restart it(a restart is not externally visible). In case of a restart, the child Actor’s preRestart() gets called. A new instance of the actor is created, after which its postRestart() method gets called. No message gets processed between the failure and the restart
- An actor can be restarted several times
- An actor can finally be stopped. It sends Stop to the context and its postStop() method will be called

- An Actor has the following methods that can be overridden:

```
trait Actor {
  def preStart(): Unit
  // the default behavior is to stop all children
  def preRestart(reason: Throwable, message: Option[Any]): Unit
  // the default behavior is to call preStart()
  def postRestart(reason: Throwable): Unit
  def postStop(): Unit
}
```

Lifecycle Monitoring

- To remove the ambiguity where a message doesn't get a response because the recipient stopped or because the network is down, Akka supports Lifecycle Monitoring, aka **DeathWatch**:
 - * An Actor registers its interest using **context.watch(target)**
 - * It will receive a **Terminated(target)** message when the target stops
 - * It will not receive any direct message from the target thereafter
- The watcher receives a **Terminated(actor: ActorRef)** message:
 - * It is a special message that our code cannot send
 - * It comes with two implicit Boolean flags: **existenceConfirmed** (was the watch sent when the target was still existing?) and **addressTerminated** (the watched actor was detected as unreachable)
 - * Terminated message are handled by the actor context, so cannot be forwarded

The Error Kernel Pattern

- Keep important data near the root, delegate the risk to the leaves:
 - * Restarts are recursive
 - * As a result, restarts are more frequent near the leaves
 - * Avoid restarting Actors with import states

Event Stream

- Because Actors can only send messages to a known address, the EventStream allows publication of messages to an unknown audience

```
trait EventStream {
  def subscribe(subscriber: ActorRef, topic: Class[_]): Boolean
  def unsubscribe(subscriber: ActorRef, topic: Class[_]): Boolean
  def unsubscribe(subscriber: ActorRef): Unit
  def publish(event: AnyRef): Unit
}

class MyActor extends Actor {
  context.system.eventStream.subscribe(self, classOf[LogEvent])
  def receive = { case e: LogEvent => ... }
  override def postStop(): Unit = {
    context.system.eventStream.unsubscribe(self)
  }
}
```

- Unhandled messages are passed to the Actor's ***unhandled(message: Any)*** method

Persistent Actor State

- The state of an Actor can be stored on disk to prevent data loss in case of a system failure
- There are two ways for persisting state:
 - * in-place updates mimics what is stored on the disk. This solution allows a fast recovery and limits the space used on the disk
 - * persist changes in append-only fashion. This solution allows fast updates on the disk. Because changes are immutable, they can be freely replicated. Finally it allows to analyze the history of a state
- Each strategy has its upsides and downsides in terms of performance to change the state, recover the state, etc.
- The stash trait allows to buffer:

```
class MyActor extends Actor with Stash {
  var state: State = ...
  def receive = {
    case NewState(text) if !state.disabled =>
      ... // sends the event to the log
      context.become(waiting, discardOld = false)
  }
  def waiting(): Receive = {
    case e: Event =>
      state = state.updated(e) // updates the state
      context.unbecome(); // reverts to the previous behavior
      unstashAll() // processes all the stashed messages
      case _ => stash() // stashes any message while waiting
  }
}
```